

Hierarchical Bayesian Modeling with tidymodels

```
# Load packages
library(tidyverse)
library(tidymodels)
library(multilevelmod) # install from Github

# Grab data and subset to the North Rocky Forest
df <- read_csv("../data/subsets/dat_small.csv")
m333 <- df %>%
  filter(province == "M333")
```

Fitting a model

```
set.seed(37) # the best number

# HB engine, standard priors
hb_spec <- linear_reg() %>%
  set_engine("stan-glmer",
    prior_aux = rstanarm::exponential(rate = 1),
    prior = NULL,
    prior_intercept = NULL,
    prior_covariance = rstanarm::decov(shape = 2))

# Fit the model
hb_fit <-
  hb_spec %>%
  fit(BIOLIVE_TPA ~ 1 + forprob + (1 | subsection), # varying intercepts
    data = m333)

hb_fit
```

```
## parsnip model object
##
## Fit time: 47.2s
## stan_glmer
## family:      gaussian [identity]
## formula:      BIOLIVE_TPA ~ 1 + forprob + (1 | subsection)
## observations: 3003
## -----
##              Median MAD_SD
## (Intercept) -4.7      3.8
## forprob      48.8      2.9
##
## Auxiliary parameter(s):
##              Median MAD_SD
## sigma 35.2      0.5
##
## Error terms:
## Groups      Name      Std.Dev.
```

```
## subsection (Intercept) 12
## Residual 35
## Num. levels: subsection 20
##
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

```
hb_stan <- hb_fit$fit$stanfit
print(hb_stan)
```

```
## Inference for Stan model: continuous.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##               mean se_mean   sd      2.5%
## (Intercept)    -4.65     0.12  3.76    -11.88
## forprob       48.80     0.05  2.84     43.26
## b[(Intercept) subsection:M333Aa]    1.33     0.12  4.19     -6.56
## b[(Intercept) subsection:M333Ab]    2.42     0.11  4.55     -6.57
## b[(Intercept) subsection:M333Ac]   -2.16     0.11  4.22    -10.56
## b[(Intercept) subsection:M333Ad]    7.05     0.14  6.28     -4.85
## b[(Intercept) subsection:M333Ai]   -4.00     0.12  5.71    -15.24
## b[(Intercept) subsection:M333Ba]    8.29     0.12  3.57     1.61
## b[(Intercept) subsection:M333Bb]   -6.30     0.12  3.35    -12.84
## b[(Intercept) subsection:M333Bc]   -2.83     0.12  3.68    -10.15
## b[(Intercept) subsection:M333Ca]   -4.72     0.12  4.59    -13.79
## b[(Intercept) subsection:M333Cb]   -2.92     0.11  3.42     -9.63
## b[(Intercept) subsection:M333Cc]   -3.13     0.12  6.14    -15.24
## b[(Intercept) subsection:M333Ce]  -10.28     0.12  4.56    -19.16
## b[(Intercept) subsection:M333Cf]  -10.10     0.12  4.72    -19.43
## b[(Intercept) subsection:M333Cg]  -12.50     0.12  4.68    -21.69
## b[(Intercept) subsection:M333Ch]  -17.21     0.12  4.73    -26.62
## b[(Intercept) subsection:M333Da]   18.22     0.12  3.63     11.33
## b[(Intercept) subsection:M333Db]   19.45     0.12  3.62     12.63
## b[(Intercept) subsection:M333Dc]   -8.06     0.12  3.66    -15.06
## b[(Intercept) subsection:M333Dd]   14.33     0.12  3.76     7.29
## b[(Intercept) subsection:M333De]   11.62     0.12  3.83     4.26
## b[(Intercept) subsection:_NEW_subsection] -0.08     0.19 11.93    -23.81
## sigma        35.17     0.01  0.45     34.31
## Sigma[subsection:(Intercept),(Intercept)] 145.38     2.00 56.22     69.42
## mean_PPD      38.57     0.01  0.91     36.77
## log-posterior -15033.76     0.17  4.68 -15043.54
##
##               25%      50%      75%
## (Intercept)   -7.18    -4.74    -2.12
## forprob       46.86    48.83    50.73
## b[(Intercept) subsection:M333Aa]   -1.41     1.36     4.08
## b[(Intercept) subsection:M333Ab]   -0.66     2.37     5.40
## b[(Intercept) subsection:M333Ac]   -4.92    -2.17     0.63
## b[(Intercept) subsection:M333Ad]    2.77     7.01    10.93
## b[(Intercept) subsection:M333Ai]   -7.82    -3.98    -0.19
## b[(Intercept) subsection:M333Ba]    5.86     8.15    10.68
## b[(Intercept) subsection:M333Bb]   -8.50    -6.31    -4.04
## b[(Intercept) subsection:M333Bc]   -5.32    -2.80    -0.34
```

```

## b[(Intercept) subsection:M333Ca]          -7.81      -4.72      -1.63
## b[(Intercept) subsection:M333Cb]          -5.23      -2.96      -0.73
## b[(Intercept) subsection:M333Cc]          -7.30      -3.08       1.00
## b[(Intercept) subsection:M333Ce]         -13.36     -10.27      -7.19
## b[(Intercept) subsection:M333Cf]         -13.23     -10.06      -6.86
## b[(Intercept) subsection:M333Cg]         -15.52     -12.44      -9.37
## b[(Intercept) subsection:M333Ch]         -20.36     -17.11     -13.95
## b[(Intercept) subsection:M333Da]          15.82      18.19      20.63
## b[(Intercept) subsection:M333Db]          17.05      19.45      21.83
## b[(Intercept) subsection:M333Dc]         -10.52      -8.10      -5.63
## b[(Intercept) subsection:M333Dd]          11.81      14.29      16.70
## b[(Intercept) subsection:M333De]           9.07      11.57      14.08
## b[(Intercept) subsection:_NEW_subsection]  -7.86      -0.15       7.17
## sigma                                     34.87      35.17      35.48
## Sigma[subsection:(Intercept),(Intercept)] 105.18     133.77     174.41
## mean_PPD                                37.98      38.58      39.19
## log-posterior                          -15036.85 -15033.45 -15030.44
##                                     97.5% n_eff Rhat
## (Intercept)                            2.84      913 1.00
## forprob                                54.46     3845 1.00
## b[(Intercept) subsection:M333Aa]          9.30     1277 1.00
## b[(Intercept) subsection:M333Ab]         11.65     1676 1.00
## b[(Intercept) subsection:M333Ac]          6.20     1506 1.00
## b[(Intercept) subsection:M333Ad]         20.05     2147 1.00
## b[(Intercept) subsection:M333Ai]          7.43     2192 1.00
## b[(Intercept) subsection:M333Ba]         15.55      958 1.00
## b[(Intercept) subsection:M333Bb]          0.23      798 1.00
## b[(Intercept) subsection:M333Bc]          4.20      952 1.00
## b[(Intercept) subsection:M333Ca]          4.10     1499 1.00
## b[(Intercept) subsection:M333Cb]          3.93      898 1.00
## b[(Intercept) subsection:M333Cc]          8.97     2643 1.00
## b[(Intercept) subsection:M333Ce]         -1.41     1520 1.00
## b[(Intercept) subsection:M333Cf]         -1.13     1500 1.00
## b[(Intercept) subsection:M333Cg]         -3.50     1501 1.00
## b[(Intercept) subsection:M333Ch]         -8.32     1565 1.00
## b[(Intercept) subsection:M333Da]         25.31      968 1.00
## b[(Intercept) subsection:M333Db]         26.68      966 1.00
## b[(Intercept) subsection:M333Dc]         -0.84      884 1.00
## b[(Intercept) subsection:M333Dd]         21.98     1041 1.00
## b[(Intercept) subsection:M333De]         19.10     1101 1.00
## b[(Intercept) subsection:_NEW_subsection] 24.30     3872 1.00
## sigma                                    36.09     4129 1.00
## Sigma[subsection:(Intercept),(Intercept)] 285.27      788 1.01
## mean_PPD                                40.33     3991 1.00
## log-posterior                          -15025.48    760 1.01
##
## Samples were drawn using NUTS(diag_e) at Sun Oct 25 09:47:20 2020.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

```

```

# Let's plot the means
hb_fit_df <- data.frame(
  fitted = hb_fit$fit$fitted.values,

```

```
true = m333$BIOLIVE_TPA,
subsection = m333$subsection
)
```

```
# Load sds
conf <- read_csv("sd_df.csv")
```

```
## Warning: Missing column names filled in: 'X1' [1]
```

In this plot, I would like error bars on the HB estimate. however, since the model fits to the plot level and then I summarized these means, I am not sure what a true error bar would look like.

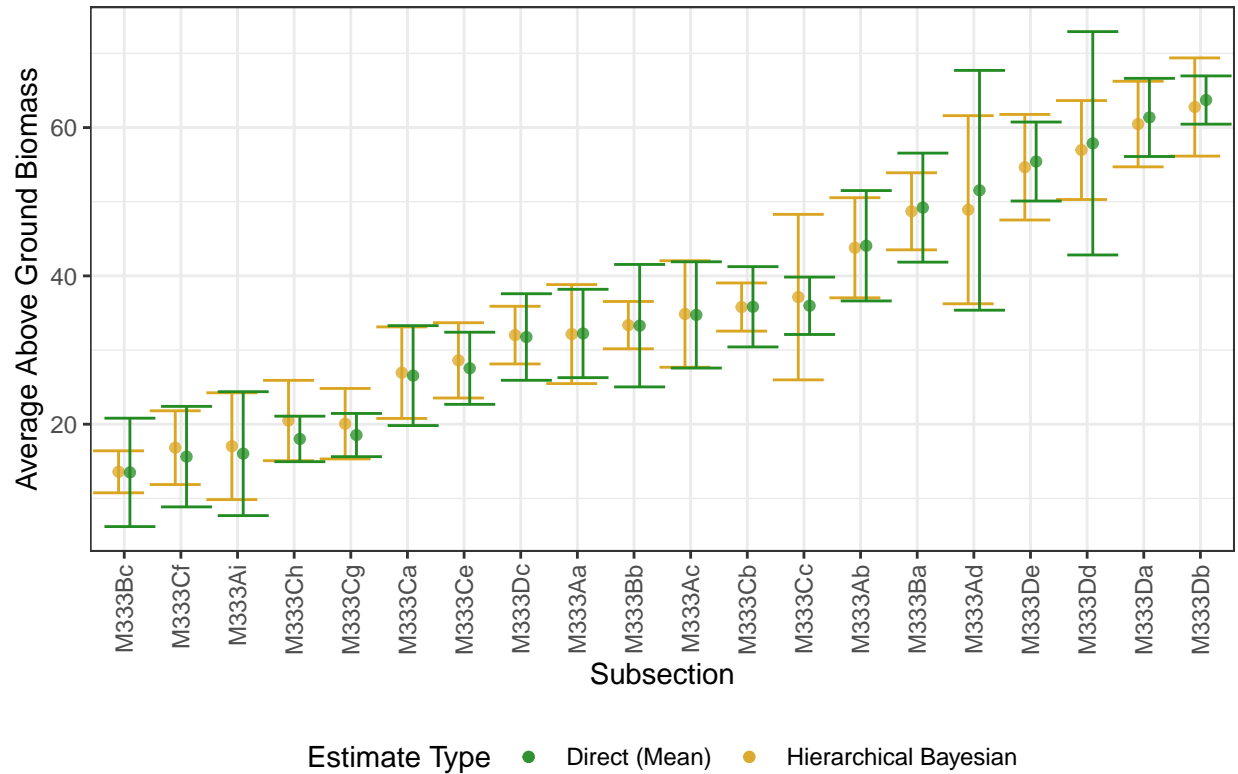
```
hb_fit_df %>%
  group_by(subsection) %>%
  summarize(mean_fit = mean(fitted),
             mean_true = mean(true)) %>%
  left_join(conf, by = c("subsection" = "id")) %>%
  mutate(subsection = fct_reorder(subsection, mean_true)) %>%
  ggplot(aes(x = subsection,
             y = mean_fit)) +
  geom_point(aes(color = "goldenrod"),
             alpha = 0.75,
             position = position_nudge(x = -0.1)) +
  geom_point(
    aes(y = mean_true, color = "forestgreen"),
    alpha = 0.75,
    position = position_nudge(x = 0.1)
  ) +
  geom_errorbar(
    mapping = aes(
      ymin = mean_fit - 1.96 * bootstrap_sd,
      ymax = mean_fit + 1.96 * bootstrap_sd
    ),
    position = position_nudge(x = -0.1),
    color = "goldenrod"
  ) +
  geom_errorbar(
    mapping = aes(
      ymin = mean_true - 1.96 * direct_sd,
      ymax = mean_true + 1.96 * direct_sd
    ),
    position = position_nudge(x = 0.1),
    color = "forestgreen"
  ) +
  theme_bw() +
  theme(axis.text.x = element_text(
    angle = 90,
    vjust = 0.5,
    hjust = 1
  ),
        legend.position = "bottom") +
  labs(
    x = "Subsection",
```

```

y = "Average Above Ground Biomass",
title = "Above Ground Biomass Estimates in the North Rocky Forest"
) +
scale_color_manual(name = 'Estimate Type',
  values = c('goldenrod'='goldenrod',
             'forestgreen'='forestgreen'),
  labels = c('Direct (Mean)', 'Hierarchical Bayesian'),
  guide = "legend")

```

Above Ground Biomass Estimates in the North Rocky Forest



```

# m333 %>%
#   group_by(subsection) %>%
#   arrange(subsection) %>%
#   summarize(n = n()) %>% View()

```

Bootstrap

```

set.seed(13)
m333_nested <- m333 %>%
  mutate(id = subsection) %>%
  group_by(subsection) %>%
  nest()
means <- data.frame()
for (j in 1:length(unique(m333$subsection))) {

```

```

for (i in 1:1000) {
  means[j, i] <- sample_n(
    m333_nested[[2]][[j]],
    size = length(m333_nested[[2]][[j]]$BIOLIVE_TPA),
    replace = TRUE
  ) %>%
  summarize(mean = mean(BIOLIVE_TPA)) %>%
  select(mean)
}
}
library(matrixStats)
sds <- means %>%
  as.matrix() %>%
  rowSds()

ordered_subsections <- m333 %>%
  arrange(subsection) %>%
  select(subsection) %>%
  unique() %>%
  unlist()

set.seed(utf8ToInt("this is gonna take awhile"))
boots <- list()
fit <- list()
mean_df <- list()
final <- data.frame()

for(i in 1:1000){
for(j in 1:length(unique(m333$subsection))) {
  boots[[j]] <- sample_n(
    m333_nested[[2]][[j]],
    size = length(m333_nested[[2]][[j]]$BIOLIVE_TPA),
    replace = TRUE
  )
  boots_df <- bind_rows(boots)
}

fit[[i]] <- hb_spec %>%
  fit(BIOLIVE_TPA ~ 1 + forprob + (1 | id),
    data = boots_df)

mean_df[[i]] <- data.frame(fitted = fit[[i]]$fit$fitted.values,
  subsection = boots_df$id)

final[1:20, i] <- mean_df[[i]] %>%
  group_by(subsection) %>%
  summarize(mean = mean(fitted)) %>%
  select(mean)

if (i %% 50 == 1) {
  print(i)
}
}

```

```

}

boot_sds <- final %>%
  as.matrix() %>%
  rowSds()

sd_df <- data.frame(
  id = ordered_subsections,
  direct_sd = sds,
  bootstrap_sd = boot_sds
)

saveRDS(sd_df, file = "sd_df.rds")

```

Fit freq model, split data, compare test MSEs

```

set.seed(1)
m333_test <- m333 %>%
  sample_frac(0.25)
m333_train <- m333 %>%
  anti_join(m333_test)

freq_spec <- linear_reg() %>%
  set_engine("lmer")

freq_fit <- freq_spec %>%
  fit(BIOLIVE_TPA ~ 1 + forprob + (1 | subsection), # varying intercepts
    data = m333_train)

hb_fit_train <-
  hb_spec %>%
  fit(BIOLIVE_TPA ~ 1 + forprob + (1 | subsection), # varying intercepts
    data = m333_train)

results_test <- hb_fit_train %>%
  predict(new_data = m333_test) %>%
  mutate(
    truth = m333_test$BIOLIVE_TPA,
    model = "hb"
  ) %>%
  bind_rows(
    freq_fit %>%
    predict(new_data = m333_test) %>%
    mutate(
      truth = m333_test$BIOLIVE_TPA,
      model = "freq"
    )
  )

results_test %>%
  group_by(model) %>%

```

```
rmse(truth = truth,  
      estimate = .pred)
```

```
## # A tibble: 2 x 4  
##   model .metric .estimator .estimate  
##   <chr> <chr>   <chr>         <dbl>  
## 1 freq  rmse     standard      34.6  
## 2 hb    rmse     standard      34.6
```

```
# tidyposterior::perf_mod() will likely be helpful comparing models
```